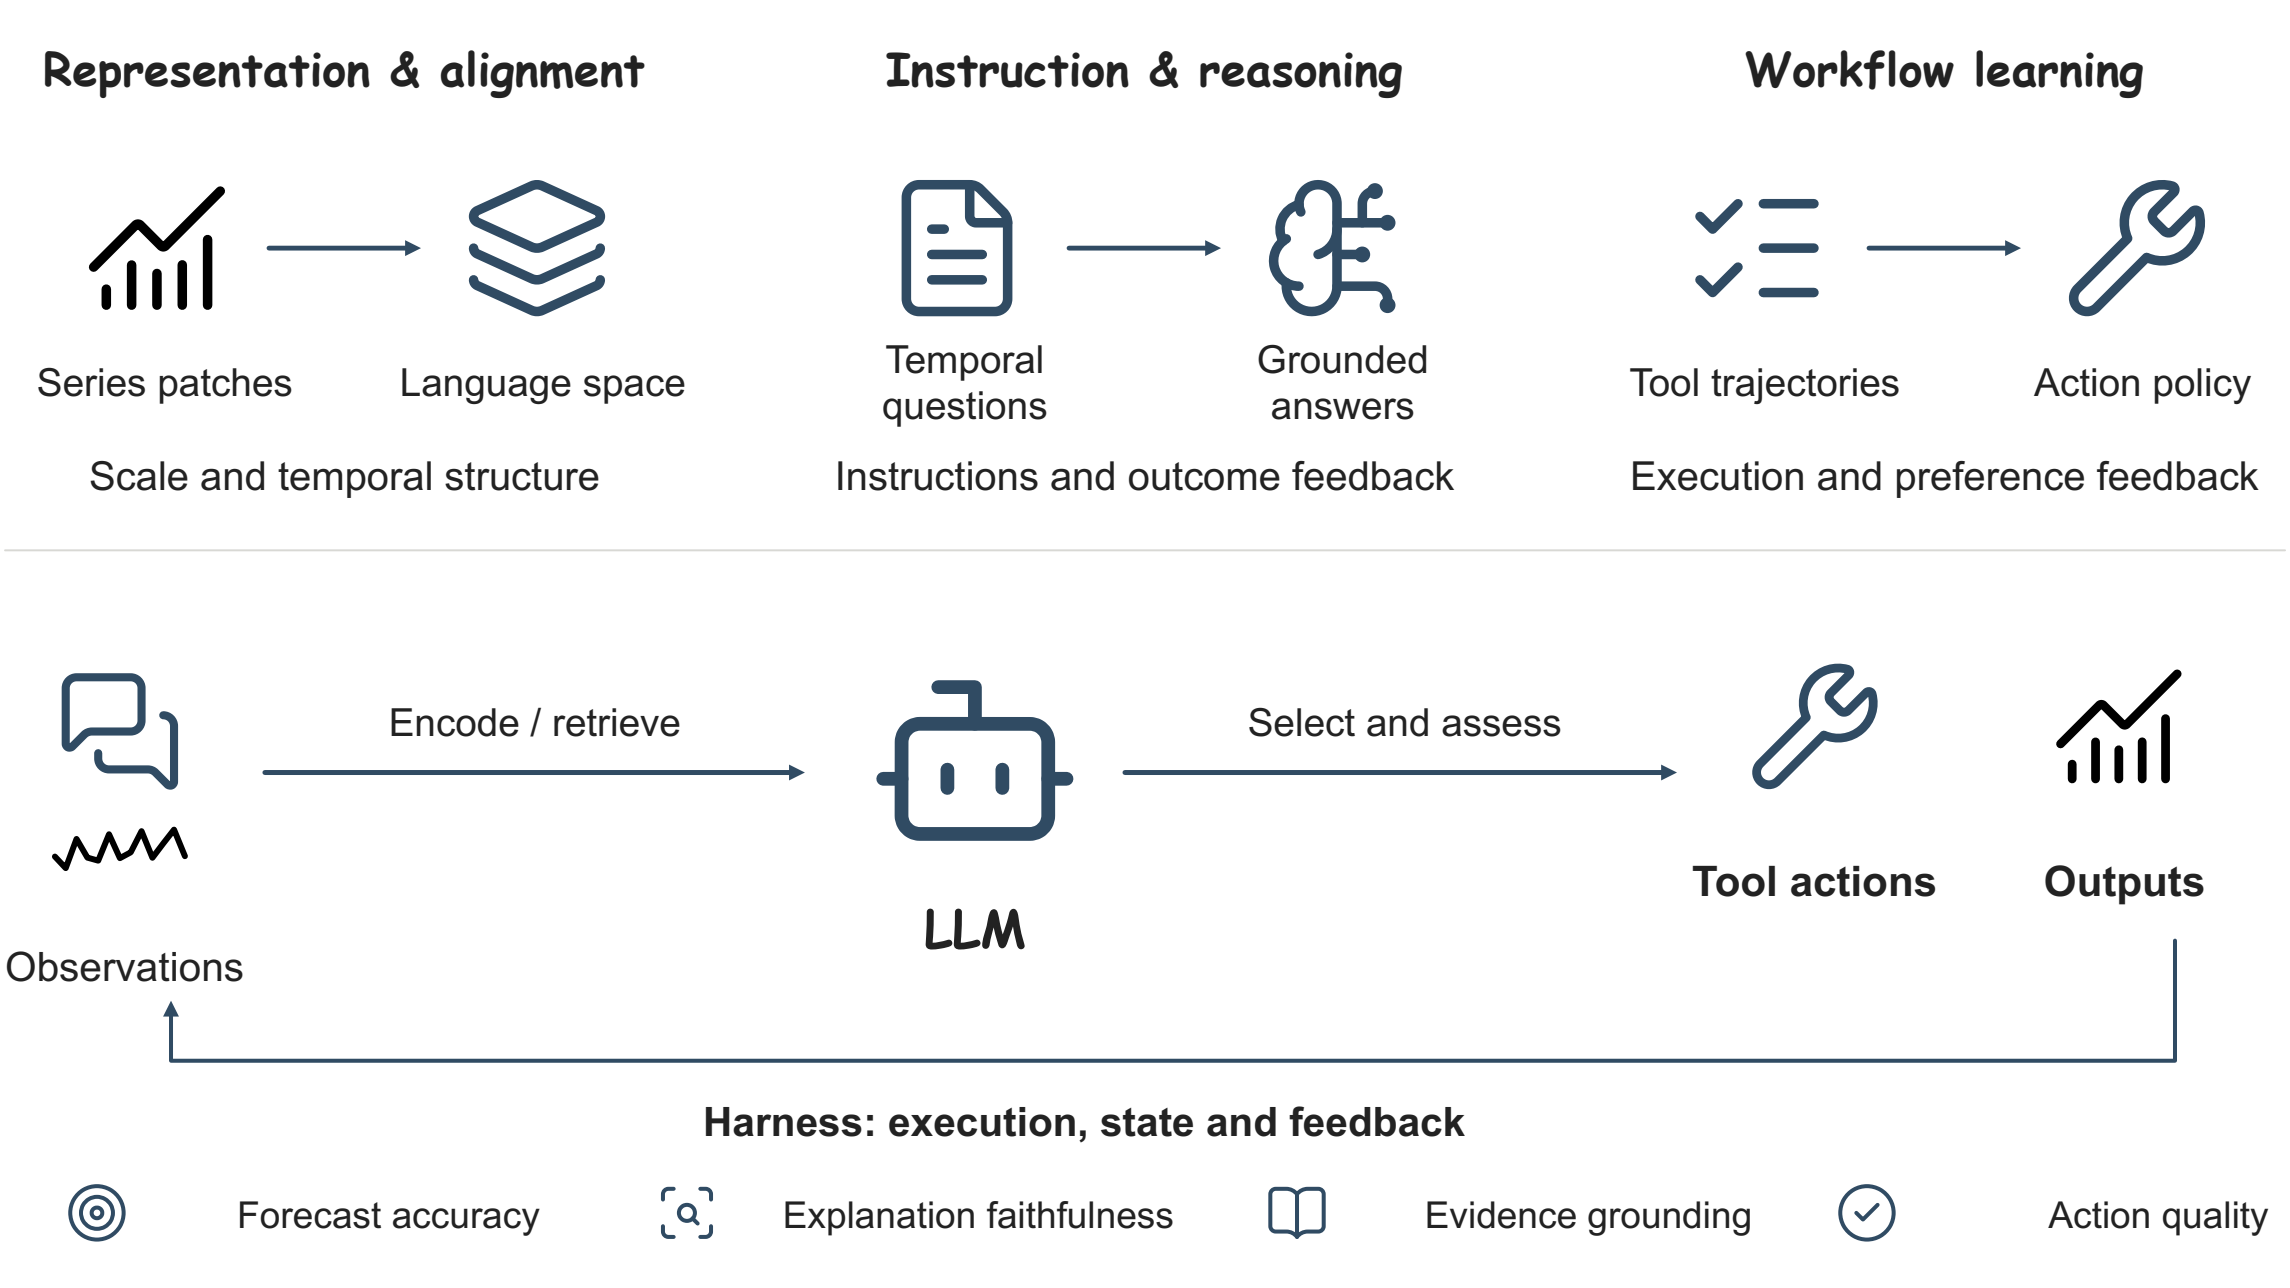
Representation & alignment
Instruction & reasoning
Workflow learning
Series patches
Language space
Temporal questions
Grounded answers
Tool trajectories
Action policy
Scale and temporal structure
Instructions and outcome feedback
Execution and preference feedback
Encode / retrieve
Select and assess
Observations
LLM
Tool actions
Outputs
Harness: execution, state and feedback
Forecast accuracy
Explanation faithfulness
Evidence grounding
Action quality